

A black and white photograph of a cemetery. In the foreground, there are several tombstones of various shapes and sizes. In the background, a tall palm tree stands on the right side. The sky is dark, and a bright, disc-shaped object, resembling a UFO, is visible in the upper center. The title "Optimal Sequence Alignment" is overlaid in a large, elegant, italicized serif font across the middle of the image.


Optimal Sequence Alignment

Overview

- The alignment problem
- The dynamic programming solution
- Pairwise alignment: exact global and local solutions
- Multiple alignment and the cost of perfection

Recap: protein scoring

C matrix – scaled frequencies of change from amino acid a to amino acid b (based on observed changes in some set)


$$\left(\frac{C_{a,b}}{f(a)f(b)} \right)$$

Expectation based solely on frequencies of amino acids (changes not favoured / disfavoured)

Better than random: ratio > 1

Random: ratio $= 1$

Worse than random: ratio < 1

log transformation

$$D_{a,b} = \boxed{S} \cdot \boxed{\log} \left(\frac{C_{a,b}}{f(a)f(b)} \right)$$

Magic  

Better than random: $D_{a,b} > 0$

Random: $D_{a,b} = 0$

Worse than random: $D_{a,b} < 0$

PAM150 matrix (S = 2, log base 2)
Half-bits

	C	S	T	P	A	G	N	D	E	Q	H	R	K	M	I	L	V	F	Y	W	
C	9																				C
S	-1	4																			S
T	-1	1	5																		T
P	-3	-1	-1	7																	P
A	0	1	0	-1	4																A
G	-3	0	-2	-2	0	6															G
N	-3	1	0	-2	-2	0	6														N
D	-3	0	-1	-1	-2	-1	1	6													D
E	-4	0	-1	-1	-1	-2	0	2	5												E
Q	-3	0	-1	-1	-1	-2	0	0	2	5											Q
H	-3	-1	-2	-2	-2	-2	1	-1	0	0	8										H
R	-3	-1	-1	-2	-1	-2	0	-2	0	1	0	5									R
K	-3	0	-1	-1	-1	-2	0	-1	1	1	-1	2	5								K
M	-1	-1	-1	-2	-1	-3	-2	-3	-2	0	-2	-1	-1	5							M
I	-1	-2	-1	-3	-1	-4	-3	-3	-3	-3	-3	-3	-3	1	4						I
L	-1	-2	-1	-3	-1	-4	-3	-4	-3	-2	-3	-2	-2	2	2	4					L
V	-1	-2	0	-2	0	-3	-3	-3	-2	-2	-3	-3	-2	1	3	1	4				V
F	-2	-2	-2	-4	-2	-3	-3	-3	-3	-3	-1	-3	-3	0	0	0	-1	6			F
Y	-2	-2	-2	-3	-2	-3	-2	-3	-2	-1	2	-2	-2	-1	-1	-1	-1	3	7		Y
W	-2	-3	-2	-4	-3	-2	-4	-4	-3	-2	-2	-3	-3	-1	-3	-2	-3	1	2	11	W
	C	S	T	P	A	G	N	D	E	Q	H	R	K	M	I	L	V	F	Y	W	

DNA matrix

Something like this usually works:

	A	G	C	T
A	1	-1	-1	-1
G	-1	1	-1	-1
C	-1	-1	1	-1
T	-1	-1	-1	1

Or this:

	A	G	C	T
A	1	0.5	-1	-1
G	0.5	1	-1	-1
C	-1	-1	1	0.5
T	-1	-1	0.5	1

Back to the alignment problem

Given a scoring scheme...

and a set of homologous sequences $S...$

introduce gaps if necessary to generate an alignment that optimizes the score

So let's make some alignments!

Sequence S_1 : length m

Sequence S_2 : length n

So let's make some alignments!

Sequence S_1 : length m

Sequence S_2 : length n

In total, there are $\binom{n+m}{m}$ possible alignments of these sequences

$n = m = 2$:

$4!/2!2! = 6$ possibilities

AB--	AB-	AB-	AB	A-B	-AB
--CD	-CD	C-D	CD	-CD	CD-

$n = m = 10$: 184,756 possible alignments

Alignment of 2 sequences, each 100 amino acids in length:

= $9.05485147 \times 10^{58}$ possibilities

Brute force is **not** going to work here...

Scaling of algorithms: Big-O Notation

What rate do resources (time, memory) increase as the input increases?

Asymptotic: upper bound on growth as input tends to infinity


Keep only dominant term:

$$3n^2 + 7n + 42 \text{ is } O(n^2)$$

Brute-force sequence alignment:


$$2n \text{ choose } n \sim O(4^n)$$

Big O complexity



The Key to Alignment

If we were given a **point X** within an optimal alignment of S_1 and S_2 , we could **split on X** and solve each problem independently



But we **don't know** any X, so divide and conquer isn't going to work

However...


In searching for the best alignment:

- Start at the beginning of the sequences and consider **every possible X**

- BUT -

- Store only the **best path** (series of matches and gaps) that leads us to X

Consider an alignment search for **AWGHE** vs **AWHEA**:



(these continue as well)

= Dynamic Programming


Consider an alignment of AWGHE vs AWHEA:

Sequence 1
----->

Sequence 2
- - - - ->

	A	W	G	H	E
A					
W					
H					
E					
A					

Each cell in the grid represents a point in the alignment where the corresponding residues have been added to the alignment



Determine the best score for every possible X from the two sequences

AWGHE vs. AWHEA	A	W	G	H	E
A	Best → (A,A)				
W					
H				Best → (H,H)	
E					
A					Best → (A,E)

Filling the matrix

We need our substitution matrix S and gap penalty scheme G

(we'll start with a linear gap penalty $G = -gd$)

For each possible X , consider the three immediate precursors

Upper left-hand corner: set to 0

S = PAM250
g = 5

AWGHE vs. AWHEA		A	W	G	H	E
	0					
A						
W						
H						
E						
A						

Filling (A,A): What is the best path to get there?

S = PAM250
g = 5

AWGHE vs. AWHEA		A	W	G	H	E
	0	-5	-10	-15	-20	-25
A	-5					
W	-10					
H	-15					
E	-20					
A	-25					

insert gap in AWGHE

insert gap in AWHEA

match

$\overline{A}$
A
A
 $\overline{A}$
A
A

Choosing the best path to (A,A)

$$S(A,A) = 2$$

Therefore:

Insert -10

Insert -10

Match 2

AWGHE vs. AWHEA	A W G H E					
	A	W	G	H	E	
A	0	-5	-10	-15	-20	-25
W	-5	2				
H	-10					
E	-15					
A	-20					
	-25					

General form: best path to any matrix cell

AWGHE vs. AWHEA							
		A	W	G	H	E	
		0	-5	-10	-15	-20	-25
A	-5						
W	-10						
H	-15						
E	-20						
A	-25						

F(2,2)

F(2,3)

F(3,2)

F(3,3)

= ?

</

Remember
paths INTO
(not out of)
each cell

AWGHE vs. AWHEA						
	A	W	G	H	E	
	0	-5	-10	-15	-20	-25
A	-5	2 → -3 → -8 → -13 → -18				
W	-10	-3	19 → 14 → 9 → 4			
H	-15	-8	14	17	20 → 15	
E	-20	-13	9	14	18	24
A	-25	-18	4	10	13	19

Global Exact Alignment: Needleman-Wunsch

Since we have retained the best path to each $F(x,y)$ in the matrix, we can trace back from the endpoint $F(m,n)$ to the origin and retrieve the optimal alignment path

Traceback for the optimal alignment

AWGHE

vs.


AWHEA

AWGHE-
AW-HEA

		A	W	G	H	E	
		0	-5	-10	-15	-20	-25
A		-5	2	-3	-8	-13	-18
W		-10	-3	19	14	9	4
H		-15	-8	14	17	20	15
E		-20	-13	9	14	18	24
A		-25	-18	4	10	13	19

Local Exact Alignment: Smith-Waterman

- Only return 'good' sub-alignments of the whole problem
- Useful, for instance, when



One more rule for local alignment

AWGHE vs. AWHEA	A	W	G	H	E
	0	0	0	0	0
A	0				
W	0				
H	0				
E					
A					

$F(2,2)$ $F(2,3)$
 
 $F(3,2)$ $F(3,3) = ?$


$F(3,3) = \max \left[\begin{array}{l} F(2,2) + S(G,H) \\ F(2,3) - d \\ F(3,2) - d \end{array} \right]$

match

insert gap in AWGHE

insert gap in AWHEA

Nothing is >0

The Needleman-Wunsch Matrix, Again

AWGHE-
AW-HEA

AWGHE vs. AWHEA		A	W	G	H	E
A	0	-5	-10	-15	-20	-25
W	-5	2	-3	-8	-13	-18
H	-10	-3	19	14	9	4
E	-15	-8	14	17	20	15
A	-20	-13	9	14	18	24
A	-25	-18	4	10	13	19

The Smith-Waterman matrix

Slightly modified
(non-trivial) S-W
example


Find the **largest**
value in the matrix,
and trace back from
there to 0

HE

HE

AWGHE vs. AYHEA						
	A	W	G	H	E	
	0	0	0	0	0	0
A	0	2	0	1	0	0
Y	0	0	2	0	0	0
H	0	0	0	0	6	1
E	0	0	0	0	1	10
A	0	2	0	1	0	5

Affine Gap Penalties



Opening a new gap
(cost = d)

Extending a gap
(cost = e)

A horizontal move now has two possible costs; we need to consider **both alternatives**

(and therefore store the best scores for each box given horizontal, vertical, or diagonal entry)

Significance of S-W Alignments

Permutation test: Randomize the alignment n times, compute mean and standard deviation

Compute **Z-score** for each replicate:

The diagram shows the formula for the Z-score of an S-W alignment, $Z(A, B)$, with three components highlighted in pink boxes and annotated with arrows:


- True score:** An arrow points to the $S(A, B)$ term in the numerator.
- Mean of replicates:** An arrow points to the $\tilde{m}$ term in the numerator.
- Standard deviation of replicates:** An arrow points to the $\tilde{\sigma}$ term in the denominator.

$$Z(A, B) = \frac{S(A, B) - \tilde{m}}{\tilde{\sigma}}$$

Significance of S-W Alignments


$$Z(A,B) = \frac{S(A,B) - \tilde{m}}{\tilde{\sigma}}$$

Curve = null model
of Z-score fit to
Gumbel
extreme value
distribution



Statistically “significant”
alignments (small p-value)

Alignment Significance



Small Z-score, large
p-value: poor alignment,
possibly not homologous?

Large Z-score, small
p-value: statistically
significant, probably
homologous


Alignment Complexity

- For each possible matching of a residue from sequence S_1 with a residue from S_2 , we need to carry out a constant number of computations and comparisons


- Total = 3 x m x n = $O(mn)$
- $\sim O(n^2)$ if we assume $m \cong n$
- Quadratic scaling!

	A	W	G	H	E
A					
W		$F(2,2)$	$F(2,3)$		
H			$F(3,2)$	$F(3,3)$ = ?	
E					
A					

Big O complexity




Multiple Sequence Alignment



Multiple Sequence Alignment

- Dynamic programming on k sequences, each of length n requires construction of a k -dimensional matrix with n^k entries

- $= O(n^k)$



- Therefore **exponential** in the number of sequences!

Scoring MSAs

- In pairwise alignment, we are optimizing the score between two sequences
- When aligning 3 or more sequences, instead optimize the **sum of pairs score**:

Sequence 1
Sequence 2
Sequence 3
Sequence 4

N
Q
Q
D

$$SP(N, Q, Q, D) =$$


$$\begin{aligned} &S(N, Q) \\ &+ S(N, Q) \\ &+ S(N, D) \\ &+ S(Q, Q) \\ &+ S(Q, D) \\ &+ S(Q, D) \end{aligned}$$

A Key Principle of MSAs

- Aligning everything at once using exact DP = n^k :
BAD
- Aligning pairs of sequences using exact DP, then
doing something with this information = $\binom{k}{2}(n^2)$:
acceptable

An Important Observation

The best alignment between a pair of sequences may not appear in the optimal multiple alignment



And

- The score of the optimal multiple alignment $S(a)$ can be no greater than the sum of optimal pairwise alignments $S(\hat{a}^{kl})$

$$\sum_{k < l} S(a^{kl}) \leq \sum_{k < l} S(\hat{a}^{kl})$$

- In general, the multiple alignment score will be **less than** the sum of all pairwise alignments

But how much less???

MSA (Carrillo and Lipman, 1988)

- If we can establish a lower bound σ on the multiple alignment score, then we constrain each $S(a^{kl})$:


$$S(\hat{a}^{kl}) - S(a^{kl}) \leq \sum_{k' < l'} S(\hat{a}^{k'l'}) - \sigma$$

Remember: sum of all best possible pairwise alignments!

σ high: $S(a^{kl})$ must be close to $S(\hat{a}^{kl})$
Sets a **bound** on “how much less”


The Consequence of σ

- We can compute $S(\hat{a}^{k'l'})$ for each **pair** of sequences, and fill the DP grid
- Any cell of the DP grid that gives $S(a^{kl})$ less than σ can be **discarded**




Constrain each pairwise alignment to score no less than

$$\sigma + S(\hat{a}^{kl}) - \sum_{k' < l'} S(\hat{a}^{k'l'})$$



The Last Step

- Multidimensional dynamic programming, restricted to “acceptable” band



- Still $O(n^k)$, but hopefully faster!

So we need all optimal pairwise alignments (again, way cheaper than naïve MSA)

We also need σ . How can we find it?

σ too **large**: we don't effectively constrain the search space!

σ too **small**: we may not find an optimal alignment!

Summary

- Dynamic programming allows the calculation of optimal pairwise alignments (for a given scoring scheme!)
- As soon as we go from 2 to >2 sequences, the exponential time complexity of the algorithm makes it impractical
- Need heuristics!